

Umeå University
Department of Computing Science

Object-Oriented Programming Methodology 7.5 p
5DV133

OU4 Sensor Network

Submitted
2016-05-23

Authors:

Johan Eklund (kv03jed@cs.umu.se)
Tommie Lindberg (c15t1g@cs.umu.se)
Jakob Lundin (c14jln@cs.umu.se)
Lorenz Gerber (dv15lgr@cs.umu.se, lozger03@student.umu.se)

Instructors:

Anders Broberg
Niklas Fries
Adam Dahlgren
Jonathan Westin
Erik Moström
Alexander Sutherland

Contents

1	Introduction	1
2	Compiling and Running of the Program	1
2.1	Javadoc	2
3	General Program Structure	2
4	Specific Design Decisions	2
4.1	A generic Environment	2
4.2	Time managment	3
4.3	Events - copy or clone?	3
4.4	Managing the Routing Tables	3
4.5	Where to go - nextNode mode	3
5	Limitations and Future Development	4
5.1	General Datastrcuture	4
5.2	Validation of Result	4
6	Testing Framework	4
7	Reflections and Account for Individual Contributions	4
7.1	Johan Eklund	4
7.2	Tommie Lindberg	6
7.3	Jakob Lundin	6
7.4	Lorenz Gerber	6
	References	6

1 Introduction

In predefined groups of students we were asked to create a simulation of a network consisting of a number of *nodes* with a short communication range [2]. The simulation consists of time steps, for every time step the network goes through the network and node and makes sure it does its job.

Using the *rumour routing* algorithm [1] you start with a network of *nodes*, in the network there are *messages*, signals, that travel the network to collect and to share information. These messages have a set amount of steps they can make before they vanish. The nodes detect unique *events* and *queries* can be made in any node for that event, the network will then send the *query* to find the *node* that detected the *event*.

Communication is done by using the rumour-routing algorithm [1]. An event in the simulation is an ID, a time stamp and a position of where it happened. Nodes know its neighbors and they are static in the grid. Nodes can have a maximum of 8 neighbors depending on location. Nodes have a table of events and know how far away every other node is. Nodes generate messages in the form of agents and queries. Queries should come back to the node within 45 steps if not the node should send another one. Queries are messages that carry the information of an event it's trying to find and a stack of nodes it has traveled through. The node the query passes through checks its table to see if it knows about the event, if it finds it it will send the query on its way to the event node. If it doesn't find it, it will send it to a random neighbor. If a query finds its target node it will get sent back to its original node the same way it got to the event so it can report about it.

In our simulation a query will vanish if it can't find the event in 45 steps. Agents are messages that carry a table of events and the path to these events. The agents travel through the network and synchronize with the nodes it passes. The events with the shortest distance gets saved in both the nodes and the agents tables. In our simulation the agents will vanish in 50 steps. In our simulation agents have a 50% chance to be generated when an event is detected. In our simulation the network is 50 by 50 and runs for 10000 time steps. Every 400 time steps four randomly chosen nodes are supposed to generate queries.

2 Compiling and Running of the Program

The program is written according to Java 1.7 and compiles on the commandline with standard command `javac RumourRoutingApp.java`. Invoking the compiled program is done by calling `java ./RumourRoutingApp.class` from the command line. A typical run will result in screen output similar to the following:

```
Event at x: 30 y 120, at time 1135, id 317
Event at x: 210 y 20, at time 1040, id 286
Event at x: 0 y 50, at time 1422, id 389
Event at x: 150 y 250, at time 1410, id 385
Event at x: 130 y 0, at time 1826, id 490
Event at x: 130 y 220, at time 3002, id 748
Event at x: 50 y 270, at time 4546, id 1164
Event at x: 130 y 210, at time 3848, id 977
Experiment finished
```

A number of parameters can be modified in the `RumourRoutingApp.java` file. The meaning of these will be described in the section 'Description of Program Structure'.

```

int NODES_X          = 50;
int NODES_Y          = 50;
int NO_NODES         = NODES_X * NODES_Y;
double NEW_EVENTS     = 0.0001;
int NODE_RANGE       = 15;
double PROB_AGENT     = 0.5;
int TTL_AGENT        = 10;
int TTL_QUERY        = 50;
int QUERY_NODES      = 4;
int QUERY_PERIODICITY = 400;
int TIMESTEPS        = 10000;
int NUMBER_OF_RECENT_NODES = 5;
int QUERY_RESEND_WAIT = TTL_QUERY * 8;

```

2.1 Javadoc

JavaDoc pages were created with the built-in functions in IntelliJ and can be found in the javadoc subdirectory. The pages are in HTML format and can be viewed by opening index.html in the javadoc directory with a web browser.

3 General Program Structure

The design of our rumour routing implementation from ‘OU3’ assignment proofed to be mostly feasible for implementation. Certainly, some details that were left open in the design slightly changed other structures, but mostly the presented code can be seen as an implementation of ‘OU3’. Hence below follows just a rather short account for the general design and we will focus more on practical implementation details.

A general aim of object oriented design is to develop a model where real world entities have their counterparts modelled as classes/objects with similar properties, behaviour and relations. For the current implementation, the real world entities of interest are sensor *nodes* with the capability to detect events and send/receive information. These three activities define two other abstract real world entities: *events* and *messages*. The former are *detected* and the latter can be *sent* and *received*. Finally, as we are interested to investigate specific situations involving a large number of the above mentioned entities, we also consider the *environment* a required entity in our implementation. Hence, *environment*, *node*, *event* and *message* were chosen as base classes. To modell the behaviour described as *rumour routing* algorithm, we designed specialications of the class *message*: *agent* and *query*.

4 Specific Design Decisions

4.1 A generic Environment

To obtain a generic environment that allows positioning nodes in which ever way desired, a coordinate based system was implemented. To simplify the main loop, a data structure containing the neighbour nodes in send/receive range is generated prior to the actual experiment. During the experiment, no more distances need to be calculated, only the initialized node neighbour data structure can be used.

4.2 Time management

In our simulation, time is semi continuous, turn/step based. In each time step, each node has at least once focus to do one activity. According to the assignment instruction, activities that consume such a time slot at the node are *sending* and *receiving* a message while *detecting* an event and *creating* messages does not consume a time step.

Busy States

The steps that consume a time slot at the node (sending/receiving) always involve two nodes. Hence, two scenarios related to concurrency have to be handled: (1) A node wants to send a message and the receiver node has already used its turn for the current timestep, or, (2) when the activity sequence comes to a node, that has already used his turn by receiving a message earlier in the same time step.

We defined a boolean *busy state* property on the node class which is switched to *true* after having used the current time step. In the end of a time step, the busy state of all nodes is reset back to *false*.

4.3 Events - copy or clone?

Event construction is initiated by the environment at individual nodes. The event object is first stored in a routing table at the node where it occurred. In our implementation, event objects are carried around as piece of information within messages. They have a *distance* property that indicates the distance in node jumps back to the node where the event occurred.

As event objects are carried around by messages, they have to update the distance property individually, hence, they need to be clones rather than just references pointing to the same event object. To recognize the equality of two events, the *hashCode* and *equals* methods were overridden to make use of the *eventId* property instead of the object reference.

4.4 Managing the Routing Tables

Routing tables were implemented as HashMaps with the event as Key and the *NodeId* as value. Routing tables are present both in *nodes* and in *agents*. Updating and managing the routing tables is the main action of agents. First both HashMaps are updated with events missing in one or the other and second, if an event is available at the node and at the agent, the one with the shorter distance to the event node is chosen.

Currently, these operations are done in rather clumsy nested *for loops*. The problem with the current code is that it involves accessing properties (distance to event node) of a specific key. This is not a normal supported operation of a HashMap, hence, it requires another *for loop*. A improvement to this structure will be proposed later.

4.5 Where to go - nextNode mode

Finding the next node to visit is implemented individually in *agent* and *query*. The method in the *query* class needs to handle three different cases, depending on the current status of the *query*. The method of the *search* mode is the same as the *nextNode* method of the *agent*. Both *query* and *agent* inherit from their superclass *message* a *queue* that holds the most recent visited *nodes*. In the practical implementation, instead of a *queue*, the datatype *dequeue* was chosen which allows also to check the last element added.

The *nextNode* methods of *query* for *track* and *homing* method were straightforward. The former matches his query event with the event in the routing table and read out the

nodeId. While the latter pops nodes from the route stack which holds references to the nodes travelled.

Figure 1 shows the UML diagram of the chosen design, automatically produced using IntelliJ.

5 Limitations and Future Development

5.1 General Datastructure

Current limitations are mostly two-fold: First the current implementation has a rather bad time complexity due to several nested for loops. This is a result of the chosen data structures and grouping of information. For example is currently the routing table a *HashMap* with a relation between an *Event* object and a *nodeId* (Integer). The *nodeId* represents the direction to follow for finding the event node in question. However, as event objects that, according to the *equalTo* method are identical, can contain different information (distance to event node). Hence sometimes, also *HashMap* keys need to be looked up which is not a normal operation in the interface of a Hash Map. This is the main cause for the nested *for loops*. In a improved version, event objects should also include the *nodeId* of the path to the event node as this information anyway is associated to the distance already found as property in the *Event*.

5.2 Validation of Result

Currently, the only validation that the implementation works correct has been done by elaborate debugging sessions. A help here was the very generic implementation which allowed to quickly modifying the experimental parameters. The currently available test cases asses the validity of individual functions and some smaller structures. In a future versions more general test cases should be added that unambiguously track and validate the whole lifecycles of individual objects, from instantiation to destruction. Also a number of more higher-lever situations need to be tested: What happens to a message when all neighbour nodes of the current node have already been visited? Could such a message block the message queue of a node?

6 Testing Framework

7 Reflections and Account for Individual Contributions

7.1 Johan Eklund

We had a pretty good idea what we wanted to do right from the first meeting. We made a solid UML diagram early on and we used it to make pseudo-code that helped us a lot through the implementation of the code. We got a group chat and a git setup early in the project which has helped a lot. We had an early idea of making unit tests before implementing code, in the end we did mostly the other way around, which made the tests a bit redundant. Differences in knowledge of the different tools has made some of the group members spend a lot of time learning the tools and not contribute as much as the more knowledgeable members. More meetings and joint programming may have helped. The problem specification can be read in different ways and a lot of different solutions can be made for whatever you believe the assignment to be. I assume this is part of the idea of the assignment. I made the initial CRC suggestion that later became our first UML-diagram.




Figure 1: UML diagram for implementing a sensory network application that allows testing of the rumour routing algorithm.

Worked on the position class and the main class that in the final version became Rumour-RoutingApp. Made other small contributions to other classes among them exceptions for a couple of classes. Wrote JUnit tests working with Tommie. Worked on the final report. Had some fierce battles with git.

7.2 Tommie Lindberg

We had a pretty good idea how we wanted to build the program from the start and a good pseudo code to follow. During the implementations did we discover some problems we may interrupt and had some differences how we wanted to avoid that. We discussed the possible solutions and decided the one that may work best. The co-operation of this assignment could have worked better, we work in different ways but we tried to stay in contact with each other on a mobile group message application called Slack. We did also have a common github folder that we worked with. Some people of the group didn't work so much with the assignment and one did more of it. I guess this is what we are about to learn about working together with other people on an assignment. That people work in different ways and the key is to find a way that works for the co-operation of a project. I implemented the message class in the program, we then encountered some things we wanted to change about it, so we did. My main implementation in this program are all the test classes which are built with junit tests and also added some exceptions.

7.3 Jakob Lundin

7.4 Lorenz Gerber

References

- [1] D. Braginsky and D. Estrin. Rumor routing algorithm for sensor networks. In *Proceedings of the 1st ACM international workshop on Wireless sensor networks and applications*, pages 22–31. ACM, 2002.
- [2] Umeå University, 5dv133 obligatorisk uppgift 3. <http://www8.cs.umu.se/kurser/5DV133/VT16/uppgifter/ou3/>, 2016. accessed: 2016-04-28.